OfficeBench — gemma4:31b Run Report

Snapshot 2026-10-02 17:23 · Run complete · 300 of 300 subtasks scored

Run setup

- **Model:** gemma4:31b, served via Ollama (no OpenAI anywhere, including the in-container llm app).
- **Reasoning:** the agent writes a **##Thought** before every **##Command**; the reasoning is captured per step.
- **Guardrail:** must switch to an app before using its actions, and use only the exact action names shown.
- **Budget:** up to 20 steps per subtask; evaluation uses the benchmark's own checks.

Results so far

132 passed and **168 failed** of 300 scored — **44.0%**. 0 subtasks still to run.

Agent	Single app (93)	Two apps (95)	Three apps (112)	Overall (300)
gemma4:31b (this run)	66.7% (62/93)	54.7% (52/95)	16.1% (18/112)	44.0% (132/300)
GPT-4 Omni (2024-05-13)	64.52	60.00	21.43	47.00
GPT-4 Turbo (2024-04-09)	56.99	50.63	11.61	38.00
Llama 3 70B-Instruct	39.79	41.05	5.36	27.33
Gemini 1.5 Pro (May 2024)	41.94	32.63	7.14	26.00
Qwen 2 72B-Instruct	30.23	28.42	8.04	21.16
Gemini 1.5 Flash (May 2024)	34.41	24.21	0.89	18.67
Gemini 1.0 Pro (Feb 2024)	24.73	13.68	0.89	12.33
GPT-3.5 Turbo (0125)	8.60	7.45	0.89	5.35

Leaderboard numbers are from the OfficeBench paper.

Why tasks fail

Of 168 failures so far, grouped by cause (codebook code in brackets):

Cause	Code	Count	Example tasks
Other (answer did not match the check)	A-REASON	59	1-8/0, 1-8/1, 1-8/2, 1-8/4, 1-10/4, 1-14/2 …
Stuck / invalid-action loop	A-ACTION	53	1-9/0, 1-19/1, 2-8/0, 2-11/1, 2-12/0, 2-14/0 …
Required file operation not performed	A-INTENT	21	1-3/4, 2-36/3, 2-39/1, 2-39/2, 2-41/0, 2-44/0 …
Ran out of steps (20)	A-ACTION	15	1-11/1, 3-5/0, 3-6/0, 3-12/2, 3-26/0, 3-29/0 …
Check requires an event/keyword the task does not ask for	EVAL-MISMATCH	9	1-2/0, 1-2/2, 3-6/1, 3-20/0, 3-24/0, 3-28/0 …
Wrong value written / operation not applied	A-REASON / A-INTENT	9	1-4/2, 1-4/3, 1-5/1, 1-5/2, 1-5/3, 1-6/0 …
Booked a slot that overlaps an existing event	A-STATE	1	1-2/1

Output could not be saved on Windows (harness)	C-HARNESS	1	3-52/0

Key change from the earlier no-reasoning run: the stuck / invalid-action loops that dominated before are essentially gone. Remaining failures are mostly the model computing a slightly wrong value or not applying the operation, plus a few cases where the benchmark check asks for something the task text does not state.

Example: reasoning captured per step

Task 1-1/0 (PASS): Add a meeting to Bob's calendar at 5/17/2024 10:30 a.m to 11:00 a.m

Step 0 — THINKING: To add a meeting to Bob's calendar, I need to use the calendar app. First, I must switch to the calendar app.

```
{"app": "system", "action": "switch_app", "target_app": "calendar"}
```

Step 1 — THINKING: I need to add a meeting to Bob's calendar on 5/17/2024 from 10:30 a.m. to 11:00 a.m. I will use the `create_event` action with the specified time format '%Y-%m-%d %H:%M:%S'.

```
{"app": "calendar", "action": "create_event", "user": "Bob", "summary": "Meeting",
"time_start": "2024-05-17 10:30:00", "time_end": "2024-05-17 11:00:00"}
```

Step 2 — THINKING: The task was to add a meeting to Bob's calendar at 5/17/2024 from 10:30 a.m to 11:00 a.m. Looking at the history, the event has been successfully created in Step 1. Therefore, the task is complete.

```
{"app": "system", "action": "finish_task", "answer": "None"}
```

Generated from the live run data. The interactive dashboard (progress, full per-task traces, filters) is the companion to this snapshot.